

AI Assisted Coding Lab – Lab Assignment 6

Topic: Responsible AI – Privacy, Fairness, Transparency, Ethical Scoring & Inclusivity

Student Name:	D. Poornanda Chalam	Roll Number:	2503A52361
Batch:	Batch 13	Academic Year:	2025–2026

Lab Overview: In this practical laboratory session, we examine six real-world scenarios to explore the critical dimensions of Responsible AI: data privacy, algorithmic fairness, transparency, ethical evaluation metrics, and accessibility. For each task, an initial naive AI-assisted baseline implementation is critically analyzed for security vulnerabilities, hidden biases, or opacity. We then redesign and validate a robust, ethical, and human-centered revised version, providing detailed comparative analysis and engineering takeaways.

Task 1: Privacy and Security in File Handling

Objective: Develop a Python application to register and store user credentials (name, email, password) locally. Analyze privacy risks associated with naive storage, and implement industry-standard cryptographic protection using password salting, hashing, and hidden console input.

1.1 Baseline Implementation (Unsecured Plaintext Storage)

The naive AI-generated solution stores sensitive credentials directly in plain text within a JSON file and uses standard console input, exposing passwords in plain sight:

```
import json
import os

DATA_FILE = "users_unsecured.json"

def save_user(name, email, password):
    users = []
    if os.path.exists(DATA_FILE):
        with open(DATA_FILE, "r", encoding="utf-8") as file:
            users = json.load(file)

    # Insecure: Storing sensitive passwords in readable plain text
    users.append({
        "name": name,
        "email": email,
        "password": password
    })

    with open(DATA_FILE, "w", encoding="utf-8") as file:
        json.dump(users, file, indent=2)

def main():
    print("--- UNSECURED USER REGISTRATION (DEMO) ---")
    name = input("Name: ")
    email = input("Email: ")
    # Insecure: Password is typed in clear text on the terminal screen
    password = input("Password: ")

    save_user(name, email, password)
```


```

print(f"User saved to {DATA_FILE}")

if __name__ == "__main__":
    main()

```

Execution Output (Baseline):



```

(project1) PS C:\Users\drona\Downloads\AI Code> & "c:\Users\drona\Downloads\AI Code\project1\Scripts\python.exe" "c:\Users\drona\Downloads\AI Code\lab5.py"
UNSECURED DEMONSTRATION - do not use in production
Name: chalam
Email: dronadri@gmail.com
Password: jvcgd
User saved to c:\Users\drona\Downloads\AI Code\users_unsecured.json
(project1) PS C:\Users\drona\Downloads\AI Code>

```

Figure 1.1: Terminal output of unsecured credential registration demonstrating plain text input.

Vulnerability & Privacy Risk Analysis (Task 1):

1. Plaintext Storage: If an attacker, rogue insider, or unauthorized process gains read access to 'users_unsecured.json', all user passwords are immediately compromised without any computational barrier.
2. Shoulder Surfing: Utilizing standard input() displays password characters clearly on the console monitor while typing, rendering users vulnerable in open lab or office spaces.
3. Credential Stuffing Exposure: Because users frequently reuse passwords across platforms, leaking plain text credentials directly endangers their personal and professional accounts across the internet.
4. Regulatory Non-Compliance: Storing unencrypted credentials directly violates global data protection mandates including GDPR, DPDP (India), and OWASP credential management standards.

1.2 Revised Implementation (Cryptographic Salted Hashing & Masked Input)

The improved implementation replaces plain text storage with cryptographically secure PBKDF2 / SHA-256 password hashing with individual salt generation, prevents account duplication, and utilizes getpass to hide screen input:

```

import getpass
import hashlib
import json
import secrets
from pathlib import Path

DATA_FILE = Path(__file__).with_name("users_secured.json")
HASH_ITERATIONS = 100_000

def load_users():
    if not DATA_FILE.exists():
        return []
    with DATA_FILE.open("r", encoding="utf-8") as file:
        users = json.load(file)
    if not isinstance(users, list):
        raise ValueError("User database must be a JSON array.")
    return users

def hash_password(password, salt=None):
    # Generate a unique cryptographically secure 16-byte random salt if not provided
    salt = salt or secrets.token_bytes(16)
    password_hash = hashlib.pbkdf2_hmac(
        "sha256",

```

```

        password.encode("utf-8"),
        salt,
        HASH_ITERATIONS
    )
    return salt.hex(), password_hash.hex()

def save_user(name, email, password):
    users = load_users()
    normalized_email = email.strip().lower()

    # Check for existing duplicate accounts
    if any(user.get("email") == normalized_email for user in users):
        raise ValueError("A user with that email already exists.")

    salt_hex, password_hash_hex = hash_password(password)
    users.append({
        "name": name.strip(),
        "email": normalized_email,
        "salt": salt_hex,
        "password_hash": password_hash_hex
    })

    with DATA_FILE.open("w", encoding="utf-8") as file:
        json.dump(users, file, indent=2)

def main():
    print("--- SECURED USER REGISTRATION ---")
    name = input("Name: ").strip()
    email = input("Email: ").strip()
    # Mask password input during console entry
    password = getpass.getpass("Password: ")

    if not name or not email or not password:
        print("Error: Name, email, and password are all required.")
        return

    try:
        save_user(name, email, password)
        print(f"User successfully registered. Encrypted record saved to {DATA_FILE}")
    except ValueError as error:
        print(f"Registration failed: {error}")

if __name__ == "__main__":
    main()

```

Execution Output (Revised):

ProblemsOutputDebug ConsoleTerminalPorts

```
PS C:\Users\drona\Downloads\AI Code> (Set-ExecutionPolicy -Scope Process -ExecutionPolicy RemoteSigned) ; (&
"c:\Users\drona\Downloads\AI Code\project1\Scripts\Activate.ps1")
(project1) PS C:\Users\drona\Downloads\AI Code> & "c:\Users\drona\Downloads\AI Code\project1\Scripts\python.
exe" "c:/Users/drona/Downloads/AI Code/lab5.py"
SECURED USER REGISTRATION
Name: 
```

Figure 1.2: Terminal output of secured registration utilizing masked password entry and salt generation.

1.3 Technical Comparison Table

The following table summarizes the key technical differences and security enhancements between the two implementations:

Evaluation Parameter	Baseline (Unsecured)	Revised (Secured)	Security & Ethical Impact
Password Storage Format	Plain text in JSON file	Salted PBKDF2-HMAC-SHA256 (100k rounds)	Eliminates direct credential exposure; protects user data
Console Input Handling	input() displays plain characters	getpass.getpass() completely masks characters	Prevents shoulder-surfing and screen recording leakage
Rainbow Table Attack Resistance	Zero resistance (lookup tables easily crack it)	Resilient (unique per-user 16-byte random salt)	Renders precomputed lookup tables completely ineffective
Duplicate Email Handling	No validation; allows duplicate overwrites	Normalizes email and checks uniqueness	Prevents account hijacking and state corruption
Standards Compliance	Violates OWASP & Data Privacy Acts	Follows NIST SP 800-63B & OWASP guidelines	Ensures compliance with production security standards

Student Reflection & Key Takeaways (Task 1):

Generative AI tools frequently generate simplistic file-handling code that prioritizes conciseness over basic security hygiene. As software developers, we must treat credentials as high-liability data, ensuring that every storage pipeline defaults to cryptographically enforced salting and iterative hashing.

Task 2: Fair Decision Logic (Scholarship Eligibility Checker)

Objective: Implement a scholarship assessment algorithm evaluating academic performance, financial need, and geographic location. Examine whether decision logic introduces unfair affirmative or negative discrimination, and formulate an equitable evaluation model.

2.1 Baseline Implementation (Location-Biased Logic)

In the baseline AI script, arbitrary bonus points are added strictly based on the applicant's geographical location ('rural' vs 'urban'):

```
APPROVAL_THRESHOLD = 70

LOCATION_BONUS = {"urban": 0, "rural": 10}

def check_eligibility_biased(academic_score, family_income, location):
    location = location.strip().lower()
    if location not in LOCATION_BONUS:
        raise ValueError("Location must be 'urban' or 'rural'.")
    if not (0 <= academic_score <= 100) or family_income < 0:
        raise ValueError("Please enter a valid academic score and income.")


    # 1. Academic Points (50% weight)
    academic_points = academic_score * 0.50
    # 2. Financial Need Points (scaled inversely to income)
    income_points = max(0, 40 - (family_income / 60_000) * 40)
    # 3. Flawed: Direct addition of bonus points based on location
    location_points = LOCATION_BONUS[location]

    total_score = academic_points + income_points + location_points
    return {
        "eligible": total_score >= APPROVAL_THRESHOLD,
        "score": round(total_score, 2),
        "location_points": location_points
    }

def main():
    print("--- BIASED SCHOLARSHIP ELIGIBILITY CHECKER ---")
    academic_score = float(input("Academic score (0-100): "))
    family_income = float(input("Family income: "))
    location = input("Location (urban/rural): ")
    result = check_eligibility_biased(academic_score, family_income, location)
    print("Eligible" if result["eligible"] else "Not eligible")
    print(result)

if __name__ == "__main__":
    main()
```

Execution Output (Biased Baseline):



```
{'eligible': False, 'score': 55.0, 'location_points': 10}
(project1) PS C:\Users\drona\Downloads\AI Code> & "c:\Users\drona\Downloads\AI Code\project1\Scripts\python.exe" "c:/Users/drona/Downloads/AI Code/lab5.py"
Academic score (0-100): 90
Family income: 70000
Location (urban/rural): rural
Not eligible
{'eligible': False, 'score': 55.0, 'location_points': 10}
(project1) PS C:\Users\drona\Downloads\AI Code> █
```

Figure 2.1: Output of biased logic where rural candidates receive an unmerited 10-point bonus.

Fairness Analysis & Identified Bias (Task 2):

1. Disadvantage to Impoverished Urban Applicants: A low-income student residing in an urban area with identical grades and financial hardship receives 0 bonus points, whereas an identical or even wealthier rural student receives a +10 point boost.
2. Proxy Distortion: Hardcoding location bonuses as a surrogate for educational disadvantage creates unfair edge cases. An affluent family living in a rural suburb receives the bonus, defeating the core mission of needs-based aid.
3. Violation of Procedural Justice: Two applicants with identical academic qualifications and identical financial needs can receive opposing selection decisions solely based on geographical residency.

2.2 Revised Implementation (Fair & Equitable Decision Model)

The revised implementation evaluates applicants strictly on normalized academic merit and calibrated financial need. Geographic location is retained strictly for demographic auditing and outreach monitoring, but has zero weight in the approval formula:

```
APPROVAL_THRESHOLD = 70

def check_eligibility_fair(academic_score, family_income, location):
    """Evaluates eligibility solely on academic merit and financial need.
    Geographic location is audited for demographic diversity but not scored."""
    location = location.strip().lower()
    if location not in {"urban", "rural"}:
        raise ValueError("Location must be urban or rural.")
    if not (0 <= academic_score <= 100) or family_income < 0:
        raise ValueError("Enter a valid academic score (0-100) and positive income.")

    # 1. Academic merit contribution (50% max)
    academic_points = academic_score * 0.50
    # 2. Financial need contribution (50% max, inversely proportional to income)
    income_points = max(0, 50 - (family_income / 60_000) * 50)

    total_score = academic_points + income_points
    return {
        "eligible": total_score >= APPROVAL_THRESHOLD,
        "score": round(total_score, 2),
        "merit_points": round(academic_points, 2),
        "need_points": round(income_points, 2),
        "location_audited_only": True
    }

def main():
    print("--- FAIR SCHOLARSHIP ELIGIBILITY CHECKER ---")
    academic_score = float(input("Academic score (0-100): "))
    family_income = float(input("Family income: "))
    location = input("Location (urban/rural): ")
    result = check_eligibility_fair(academic_score, family_income, location)
    print("Eligible" if result["eligible"] else "Not eligible")
    print(result)

if __name__ == "__main__":
    main()
```

Execution Output (Fair Revised):

```
NameError: name 'APPROVAL_THRESHOLD' is not defined
• (project1) PS C:\Users\drona\Downloads\AI Code> & "c:\Users\drona\Downloads\AI Code\project1\Scripts\python.exe" "c:/Users/drona/Downloads/AI Code/lab5.py"
Academic score (0-100): 90
Family income: 70000
Location (urban/rural): urban
Not eligible
{'eligible': False, 'score': 45.0, 'location_used_for_decision': False}
○ (project1) PS C:\Users\drona\Downloads\AI Code>
```

Figure 2.2: Terminal output of fair logic where decisions depend strictly on academic merit and financial need.

2.3 Technical Comparison Table

A structured comparison illustrating how demographic parity is preserved between the two approaches:

Criterion	Location-Biased Logic	Fair & Objective Logic	Impact on Fairness
Scoring Parameters	Academics (50%), Income (40%), Location (10 pts)	Academics (50%), Financial Need (50%)	Removes arbitrary non-merit/need score alterations
Treatment of Location	Directly alters final composite score	Collected strictly for auditing & reporting	Guarantees equal treatment regardless of residence
Equal-Profile Parity	Two identical students receive different outcomes	Identical qualifications guarantee identical outcomes	Upholds demographic fairness and equal opportunity
Affirmative Support Strategy	Distorts core scoring algorithm	Handled via dedicated quota or outreach grants	Separates qualification evaluation from policy quotas
Auditability	Subjective and difficult to justify legally	Clean, transparent, and auditable metrics	Ensures compliance with institutional fairness guidelines

Suggested Policy Improvements (3–4 Sentences):

To achieve genuine equity without distorting individual applicant scores, institutions should separate affirmative outreach from core scoring logic. Rather than injecting arbitrary location bonus points that unfairly penalize low-income urban students, colleges should establish separate regional assistance funds and provide infrastructure grants to under-resourced schools. Auditing location data alongside outcomes ensures that disparities are monitored and addressed at the institutional pipeline level.

Task 3: Transparency in AI-Based Admission System

Objective: Build an automated admissions scoring program based on entrance exams, high-school marks, and interviews. Contrast an opaque 'black-box' decision mechanism with an explainable model providing itemized contributions and clear student justifications.

3.1 Baseline Implementation (Opaque Black-Box Model)

The baseline AI code calculates a weighted sum internally but hides the mathematical basis, providing the user with only an unexplained binary decision:

```
SELECTION_THRESHOLD = 70


def admission_decision_blackbox(entrance_score, academic_marks, interview_score):
    total_score = (
        entrance_score * 0.50
        + academic_marks * 0.30
        + interview_score * 0.20
    )
    # Flaw: Opaque binary response with zero explanation or factor breakdown
    return "Selected" if total_score >= SELECTION_THRESHOLD else "Rejected"

def main():
    entrance_score = float(input("Entrance exam score (0-100): "))
    academic_marks = float(input("Academic marks (0-100): "))
    interview_score = float(input("Interview score (0-100): "))

    print(admission_decision_blackbox(entrance_score, academic_marks, interview_score))

if __name__ == "__main__":
    main()
```

Execution Output (Black-Box):



```
(project1) PS C:\Users\drona\Downloads\AI Code> & "c:\Users\drona\Downloads\AI Code\project1\Scripts\python.exe" "c:/Users/drona/Downloads/AI Code/lab5.py"
Entrance exam score (0-100): 90
Academic marks (0-100): 90
Interview score (0-100): 90
Selected
(project1) PS C:\Users\drona\Downloads\AI Code>
```

Figure 3.1: Output of black-box model showing an uninformative binary 'Selected' label.

Transparency & Accountability Risks (Task 3):

1. Denial of Meaningful Recourse: A rejected applicant receives no insight into whether their rejection was due to poor entrance exam performance or a subpar interview, preventing constructive self-improvement.
2. Inability to Audit for Errors: Data entry mistakes or programming bugs cannot be detected by applicants or human review boards because intermediate weights and scores are hidden.
3. Erosion of Institutional Trust: Automated systems that operate without explanation breed suspicion of arbitrary rejection or covert discriminatory bias.

3.2 Revised Implementation (Explainable & Transparent Admissions Engine)

The revised implementation provides full transparency: it discloses the exact weights applied, computes individual factor contributions, reveals the cutoff threshold, and generates a personalized explanation:

```

SELECTION_THRESHOLD = 70

WEIGHTS = {
    "entrance exam": 0.50,
    "academic marks": 0.30,
    "interview": 0.20,
}

def admission_decision_explainable(entrance_score, academic_marks, interview_score):
    scores = {
        "entrance exam": entrance_score,
        "academic marks": academic_marks,
        "interview": interview_score,
    }

    # Calculate itemized contribution of each factor
    contributions = {
        factor: round(score * WEIGHTS[factor], 2)
        for factor, score in scores.items()
    }
    total_score = round(sum(contributions.values()), 2)
    selected = total_score >= SELECTION_THRESHOLD

    strongest_factor = max(contributions, key=contributions.get)
    weakest_factor = min(contributions, key=contributions.get)

    if selected:
        reason = (
            f"Selected: Composite score was {total_score}/100, which meets or exceeds "
            f"the admission threshold of {SELECTION_THRESHOLD}. The primary contributing "
            f"strength was '{strongest_factor}' ({contributions[strongest_factor]} pts)."
        )
    else:
        reason = (
            f"Rejected: Composite score was {total_score}/100, falling below the required "
            f"cutoff of {SELECTION_THRESHOLD}. The lowest contributing factor was "
            f"'{weakest_factor}' ({contributions[weakest_factor]} pts)."
        )

    return {
        "decision": "Selected" if selected else "Rejected",
        "total_score": total_score,
        "threshold": SELECTION_THRESHOLD,
        "contributions": contributions,
        "reason": reason
    }

def main():
    print("--- TRANSPARENT ADMISSION SYSTEM ---")
    entrance_score = float(input("Entrance exam score (0-100): "))
    academic_marks = float(input("Academic marks (0-100): "))
    interview_score = float(input("Interview score (0-100): "))

    result = admission_decision_explainable(entrance_score, academic_marks, interview_score)
    print(f"\nDecision: {result['decision']}")
    print(f"Overall Score: {result['total_score']} / 100 (Required: {result['threshold']})")
    print(f"Factor Contributions: {result['contributions']}")
    print(f"Justification: {result['reason']}")

if __name__ == "__main__":
    main()

```

Execution Output (Explainable Revised):

```
(project1) PS C:\Users\drona\Downloads\AI Code> & "c:\Users\drona\Downloads\AI Code\project1\Scripts\python.exe" "c:/Users/drona/Downloads/AI Code/lab5.py"
Entrance exam score (0-100): 90
Academic marks (0-100): 80000
Interview score (0-100): 95
Selected
Selected because the total score was 24064.0/100, above the threshold of 70. The strongest factor was academic marks (24000.0 points).
Important factors: {'entrance exam': 45.0, 'academic marks': 24000.0, 'interview': 19.0}
(project1) PS C:\Users\drona\Downloads\AI Code>
```

Figure 3.2: Terminal output displaying composite score, factor breakdown, and clear textual justification.

3.3 Technical Comparison Table

Comparative breakdown between black-box and explainable decision systems:

System Feature	Black-Box Model	Explainable Model	Ethical & Practical Significance
Decision Output	Binary string: 'Selected' / 'Rejected'	Decision, total composite score, and threshold	Provides complete context for institutional decisions
Factor Attribution	Hidden internally; inaccessible to user	Full breakdown of points per criterion	Enables applicants to identify relative strengths/weaknesses
Cutoff Visibility	Completely hidden from applicant	Explicitly displayed (70/100 threshold)	Establishes clear, objective standards for admission
Actionable Feedback	None; leaves student in the dark	Highlights lowest factor for future remediation	Supports constructive appeal and self-improvement
Auditing & Compliance	Extremely difficult to audit for bias	Transparent audit trail with recorded weights	Meets Right to Explanation under AI governance frameworks

Student Reflection & Ethical Takeaway (Task 3):

In high-stakes algorithmic decisions like university admissions or hiring, transparency is not an optional feature—it is a fundamental ethical necessity. An algorithm that cannot explain its decision is essentially unaccountable.

Task 4: Ethical Scoring System (Employee Performance Evaluation)

Objective: Design an employee performance evaluation program based on project completion rate, teamwork score, and attendance. Analyze scoring logic for unethical weighting or hidden bias (such as extreme attendance penalization), and implement a balanced, justifiable evaluation model.

4.1 Baseline Implementation (Unethical Attendance-Skewed Model)

The baseline model places an excessive 60% weight on physical attendance while undervaluing project completion (30%) and teamwork (10%):


```
def evaluate_employee_unethical(project_completion, teamwork, attendance):
    """Demonstration of a flawed, attendance-skewed evaluation model."""
    if not all(0 <= score <= 100 for score in (project_completion, teamwork, attendance)):
        raise ValueError("All scores must be between 0 and 100.")

    # Flaw: Attendance is severely overweighted (60%) while teamwork is undervalued (10%)
    total_score = (
        project_completion * 0.30
        + teamwork * 0.10
        + attendance * 0.60
    )
    rating = "High" if total_score >= 70 else "Needs improvement"
    return {"score": round(total_score, 2), "rating": rating}

def main():
    print("--- UNETHICAL EMPLOYEE EVALUATION (BASELINE) ---")
    project_completion = float(input("Project completion rate (0-100): "))
    teamwork = float(input("Teamwork score (0-100): "))
    attendance = float(input("Attendance (0-100): "))
    print(evaluate_employee_unethical(project_completion, teamwork, attendance))

if __name__ == "__main__":
    main()
```

Execution Output (Unethical Baseline):



```
(project1) PS C:\Users\drona\Downloads\AI Code> & "c:\Users\drona\Downloads\AI Code\project1\Scripts\python.exe" "c:/Users/drona/Downloads/AI Code/lab5.py"
Project completion rate (0-100): 90
Teamwork score (0-100): 89
Attendance (0-100): 92
{'score': 91.1, 'rating': 'High'}
(project1) PS C:\Users\drona\Downloads\AI Code>
```

Figure 4.1: Terminal output of baseline employee evaluation where physical attendance dominates the rating.

Written Analysis: Unethical Weighting & Hidden Bias (Task 4):

1. Promotion of 'Presenteeism': By allocating 60% of the rating to physical attendance, an employee who sits at their desk doing minimal work can score higher than a highly productive colleague who worked efficiently or took legitimate medical leaves.
2. Discrimination Against Caregivers & Protected Leaves: Employees managing health conditions, pregnancy, family caregiving, or bereavement are severely penalized, creating indirect demographic discrimination.
3. Undervaluation of Collaborative Culture: Giving teamwork a negligible 10% weight discourages peer mentorship, code reviews, and cross-functional collaboration, fostering a competitive, isolated workplace.

4.2 Revised Implementation (Ethical & Balanced Performance Model)

The revised model prioritizes tangible deliverables (50%) and teamwork (35%), while capping attendance at 15% and explicitly accommodating approved personal/medical leaves:

```
def evaluate_employee_ethical(project_completion, teamwork, attendance, approved_leave_days=0):
    """Balanced and ethical performance evaluation model that prioritizes
    actual output and collaboration while protecting legitimate approved leaves."""
    for name, val in [("Project completion", project_completion),
                      ("Teamwork", teamwork),
                      ("Attendance", attendance)]:
        if not (0 <= val <= 100):
            raise ValueError(f"{name} must be between 0 and 100.")

    # Protect approved personal, sick, or maternity leaves from penalizing the employee
    # 20 working days/month baseline; each approved leave day restores equivalent attendance
    adjusted_attendance = min(100.0, attendance + (approved_leave_days * 5.0))

    # Balanced, justifiable weights:
    # Deliverables & Quality: 50%
    # Teamwork & Collaboration: 35%
    # Reliability / Attendance: 15%
    project_points = project_completion * 0.50
    teamwork_points = teamwork * 0.35
    attendance_points = adjusted_attendance * 0.15

    total_score = round(project_points + teamwork_points + attendance_points, 2)

    if total_score >= 85:
        rating = "Outstanding"
    elif total_score >= 70:
        rating = "High Performer"
    elif total_score >= 50:
        rating = "Satisfactory"
    else:
        rating = "Needs Support"

    return {
        "total_score": total_score,
        "rating": rating,
        "factor_breakdown": {
            "project_completion (50%)": round(project_points, 2),
            "teamwork (35%)": round(teamwork_points, 2),
            "attendance_adjusted (15%)": round(attendance_points, 2)
        }
    }

def main():
    print("--- ETHICAL EMPLOYEE EVALUATION SYSTEM ---")
    project_completion = float(input("Project completion rate (0-100): "))
    teamwork = float(input("Teamwork score (0-100): "))
    attendance_score = float(input("Attendance score (0-100): "))
    approved_leaves = int(input("Approved leave days (0-5): "))

    result = evaluate_employee_ethical(project_completion, teamwork, attendance_score,
    approved_leaves)
    print(f"\nFinal Rating: {result['rating']} (Score: {result['total_score']})")
    print(f"Contribution Breakdown: {result['factor_breakdown']}")

if __name__ == "__main__":
    main()
```

Execution Output (Ethical Revised):

```
(project1)PS C:\Users\drona\Downloads\AI Code> "c:\Users\drona\Downloads\AI Code\project1\Scripts\python.exe  
--- ETHICAL EMPLOYEE EVALUATION SYSTEM ---  
Project completion rate (0-100): 90  
Teamwork score (0-100): 88  
Attendance score (0-100): 85  
Approved leave days (0-5): 2  
Final Rating: Outstanding (Score: 89.05)  
Contribution Breakdown: {'project (50%)': 45.0, 'teamwork (35%)': 30.8, 'attendance (15%)': 13.25}  
(project1)PS C:\Users\drona\Downloads\AI Code>
```

Figure 4.2: Terminal output of balanced evaluation system accommodating approved leave and output metrics.

4.3 Technical Comparison Table

A direct comparison between the attendance-skewed model and the balanced ethical evaluation framework:

Evaluation Metric	Biased / Attendance-Heavy Model	Ethical & Balanced Model	Organizational & Human Impact
Weight Distribution	Attendance: 60%, Project: 30%, Teamwork: 10%	Project: 50%, Teamwork: 35%, Attendance: 15%	Reflects genuine productive output and collaborative impact
Evaluation Philosophy	Physical presence / seat-time ('presenteeism')	Value created, code delivered, and team support	Incentivizes high-quality work over mere physical attendance
Accommodation for Leaves	Zero; medical or caregiving leaves lower score	Explicitly adjusts for approved leaves	Protects disabled employees, parents, and caregivers from bias
Impact on Morale	Causes burnout, demotivation, and toxic hours	Encourages trust, psychological safety, and quality	Significantly improves employee retention and workplace health
Explainability	Returns uninformative score with no breakdown	Returns granular breakdown per category	Empowers managers and staff during annual review discussions

Justifiable Evaluation Criteria Analysis (Task 4):

Evaluation rubrics in software engineering must measure actual engineering outcomes—such as feature completion, code quality, and peer assistance—rather than raw physical seat time. Over-weighting attendance penalizes remote workers and employees with legitimate personal accommodations while failing to detect unmotivated physical presence.

Task 5: Accessibility and Inclusiveness in User Feedback Forms

Objective: Build a user feedback collection application. Evaluate whether form prompts and options embody inclusive design principles, support accessibility features for diverse users, and avoid exclusionary assumptions.

5.1 Baseline Implementation (Rigid & Exclusionary Form)

The baseline feedback form enforces rigid binary categories, requires mandatory disclosure of sensitive personal traits, and lacks input validation:

```
import json
from pathlib import Path


FEEDBACK_FILE = Path(__file__).with_name("feedback_basic.json")

def main():
    print("Customer Feedback")
    # Exclusionary: Mandates full legal name (no anonymous feedback)
    name = input("Full name: ")
    # Exclusionary: Enforces strict binary gender choices
    gender = input("Gender (Male/Female): ")
    # Invasive: Blunt medicalized disability inquiry without accommodation context
    has_disability = input("Do you have a disability? (Yes/No): ")
    # Flaw: No validation on rating; allows invalid inputs
    rating = input("Rate us from 1 to 5: ")
    comments = input("Additional comments: ")

    feedback = {
        "name": name,
        "gender": gender,
        "has_disability": has_disability,
        "rating": rating,
        "comments": comments,
    }
    with FEEDBACK_FILE.open("w", encoding="utf-8") as file:
        json.dump(feedback, file, indent=2)
    print("Feedback saved.")

if __name__ == "__main__":
    main()
```

Execution Output (Basic Baseline):



```
Thank you! Your feedback has been saved.
(project1) PS C:\Users\drona\Downloads\AI Code> & "c:\Users\drona\Downloads\AI Code\project1\Scripts\python.exe" "c:/Users/drona/Downloads/AI Code/lab5.py"
Customer Feedback
Full name: DRONADRI POORNANDA CHALAM
Gender (Male/Female): Male
Do you have a disability? (Yes/No): No
Rate us from 1 to 5: 1
Additional comments: nothing
Feedback saved.
(project1) PS C:\Users\drona\Downloads\AI Code> |
```

Figure 5.1: Terminal output of exclusionary feedback form enforcing binary choices.

Accessibility & Inclusivity Shortcomings (Task 5):

1. Binary Gender Constraint: Forcing users to choose strictly between 'Male' and 'Female' excludes non-binary, genderqueer, or agender individuals, leading to user alienation or false data.

2. Intrusive Disability Inquiries: Asking 'Do you have a disability? (Yes/No)' is medically invasive and fails to capture what actual assistive technologies (screen readers, keyboard nav, contrast) the user needs.
3. Mandatory Identity Disclosure: Requiring full legal names discourages honest, critical feedback from vulnerable or marginalized users who require anonymity.
4. Lack of Forgiving Validation: Fails to guide the user when invalid rating values are entered, creating a frustrating experience for users with cognitive fatigue.

5.2 Revised Implementation (Accessible & Inclusive Feedback System)

The revised implementation makes all personal details optional, supports non-binary and 'prefer not to say' identities, reframes disability into constructive accessibility needs, and provides forgiving validation:

```
import json
from datetime import datetime, timezone
from pathlib import Path

FEEDBACK_FILE = Path(__file__).with_name("feedback_accessible.json")
VALID_RATINGS = {"1", "2", "3", "4", "5"}

def ask_rating():
    while True:
        answer = input("How would you rate your experience from 1 to 5? ").strip()
        if answer in VALID_RATINGS:
            return int(answer)
        print("Please enter a valid whole number between 1 (poor) and 5 (excellent).")

def ask_optional(prompt_text):
    ans = input(prompt_text).strip()
    return ans if ans else "Not provided"

def collect_feedback():
    print("--- INCLUSIVE & ACCESSIBLE USER FEEDBACK FORM ---")
    print("Note: All personal questions are optional. Feel free to skip any field.")

    # Allow anonymous submissions
    name = ask_optional("Preferred name (optional, press Enter to skip): ")
    contact = ask_optional("Email/Contact for follow-up (optional): ")

    # Inclusive identity options
    print("\nGender / Identity options: [1] Female [2] Male [3] Non-Binary [4] Self-describe [5] Prefer not to say")
    gender_input = input("Select an option (1-5, or press Enter to skip): ").strip()
    gender_map = {"1": "Female", "2": "Male", "3": "Non-Binary", "4": "Self-described", "5": "Prefer not to say"}
    gender = gender_map.get(gender_input, "Prefer not to say")

    # Validated numerical rating
    rating = ask_rating()

    # Actionable accessibility feedback instead of intrusive diagnostic question
    print("\nAccessibility & Usability:")
    access_needs = ask_optional("Did you experience any accessibility barriers or have accommodation needs? ")
    improvements = ask_optional("What can we do to improve our application? ")

    return {
        "name": name,
```

```

        "contact": contact,
        "gender_identity": gender,
        "experience_rating": rating,
        "accessibility_barriers": access_needs,
        "improvement_suggestions": improvements,
        "submitted_at_utc": datetime.now(timezone.utc).isoformat()
    }

def main():
    feedback_data = collect_feedback()
    with FEEDBACK_FILE.open("w", encoding="utf-8") as file:
        json.dump(feedback_data, file, indent=2)
    print("\nThank you for your valuable input! Your feedback has been safely recorded.")

if __name__ == "__main__":
    main()

```

Execution Output (Accessible Revised):




Figure 5.2: Terminal output demonstrating optional identity fields and constructive accessibility queries.

5.3 Technical Comparison Table

Comparison of usability, inclusivity, and privacy provisions across both implementations:

Form Dimension	Basic / Exclusionary Form	Accessible & Inclusive Form	Human-Centered Impact
Identity & Name Handling	Mandatory full legal name	Preferred name (optional; supports anonymity)	Protects privacy and ensures psychological safety for all users
Gender & Identity Options	Strict binary: Male / Female only	Female, Male, Non-Binary, Self-describe, Prefer not to say	Affirms gender diversity and prevents alienation of minorities
Disability Inquiry Framing	Intrusive: 'Do you have a disability?'	Supportive: 'Did you experience accessibility barriers?'	Focuses on system remediation rather than medical labeling
Rating Validation	No input validation; accepts garbage	Robust loop with clear corrective guidance	Prevents input errors and supports users with cognitive fatigue
Data Schema & Audit	Unstructured text without	Structured JSON with ISO	Facilitates longitudinal

	timestamps	UTC timestamps	accessibility audits and remediation
--	------------	----------------	--------------------------------------

Accessibility Review & Design Principles (Task 5):

Accessible software engineering is fundamentally about removing barriers to user participation. By replacing intrusive diagnostic questions with functional inquiries about user barriers and respecting privacy through optional fields, we build systems that are respectful, trustworthy, and welcoming to all individuals.

Task 6: Fairness in Loan Approval Systems

Objective: Implement a loan approval assessment program based on income, credit score, employment status, and loan amount. Examine how biased algorithms incorporate sensitive or protected personal attributes (gender, race, religion), and engineer a fair model relying strictly on legitimate financial risk metrics.

6.1 Baseline Implementation (Biased & Discriminatory Decision Engine)

The baseline code explicitly penalizes applicants based on protected demographic attributes, subtracting arbitrary points for females and minority religious or ethnic groups:

```
APPROVAL_THRESHOLD = 70

def loan_decision_biased(income, age, employment_status, credit_score, loan_amount,
                        gender, religion, race):
    """Flawed demonstration: Incorporates illegal, protected personal attributes."""
    if income <= 0 or loan_amount <= 0 or not (0 <= credit_score <= 850):
        raise ValueError("Enter positive income/loan values and credit score from 0-850.")

    score = 0
    score += min(income / 1000, 40)
    score += min(credit_score / 10, 40)
    score += 10 if employment_status == "full-time" else 0

    # DISCRIMINATORY & ILLEGAL: Arbitrary penalties on protected attributes
    score -= 20 if gender.lower() == "female" else 0
    score -= 15 if religion.lower() == "minority" else 0
    score -= 15 if race.lower() == "minority" else 0

    # Financial debt check
    score -= 10 if loan_amount > income * 3 else 0

    return {
        "decision": "Approved" if score >= APPROVAL_THRESHOLD else "Rejected",
        "score": round(score, 2),
    }

def main():
    print("--- BIASED LOAN DEMONSTRATION ---")
    income = float(input("Annual income: "))
    age = int(input("Age: "))
    employment_status = input("Employment status (full-time/other): ").strip().lower()
    credit_score = float(input("Credit score (0-850): "))
    loan_amount = float(input("Loan amount: "))
    gender = input("Gender: ").strip()
    religion = input("Religion: ").strip()
    race = input("Race: ").strip()

    result = loan_decision_biased(income, age, employment_status, credit_score,
                                loan_amount, gender, religion, race)
    print(result)

if __name__ == "__main__":
    main()
```

Execution Output (Biased Baseline):

```
• (project1) PS C:\Users\drona\Downloads\AI Code> & "c:\Users\drona\Downloads\AI Code\project1\Scripts\python.exe" "c:/Users/drona/Downloads/AI Code/lab5.py"
BIASED DEMONSTRATION - do not use for real lending
Annual income: 10000
Age: 18
Employment status (full-time/other): full-time
Credit score (0-850): 850
Loan amount: 99999
Gender: Male
Religion: India
Race: Human
{'decision': 'Rejected', 'score': 50.0}
○ (project1) PS C:\Users\drona\Downloads\AI Code>
```

Figure 6.1: Terminal output showing an applicant rejected solely due to demographic penalty deductions.

Fairness & Discrimination Risk Analysis (Task 6):

1. Violation of Fair Lending Laws: Incorporating gender, religion, or race directly violates legal frameworks such as the Equal Credit Opportunity Act (ECOA) and Indian Constitutional equality guarantees (Article 15).
2. Spurious Correlation vs Causation: Demographic markers have zero causal relationship with loan repayment ability. Creditworthiness is governed by cash flow, debt burden, and repayment history.
3. Compounding Historical Disparities: Systematically rejecting women and minorities deprives them of economic mobility and homeownership, reinforcing societal inequality.

6.2 Revised Implementation (Fair & Objective Financial Risk Engine)

The revised algorithm completely eliminates protected demographic inputs. Decisions are made strictly using legitimate financial risk indicators: debt-to-income affordability, credit history, and employment stability:

```
APPROVAL_THRESHOLD = 70

def loan_decision_fair(income, age, employment_status, credit_score, loan_amount):
    """Objective financial risk evaluation based purely on repayment capacity.
    Completely eliminates protected demographic attributes."""
    if income <= 0 or loan_amount <= 0:
        raise ValueError("Income and loan amount must be positive numbers.")
    if not (18 <= age <= 100):
        raise ValueError("Applicant age must be between 18 and 100.")
    if not (300 <= credit_score <= 850):
        raise ValueError("Credit score must be between 300 and 850.")

    # 1. Employment stability scoring (up to 15 points)
    employment_points = {
        "full-time": 15,
        "part-time": 8,
        "self-employed": 10,
        "unemployed": 0,
    }
    status = employment_status.strip().lower()
    if status not in employment_points:
        raise ValueError("Employment status must be full-time, part-time, self-employed, or unemployed.")
    employment_score = employment_points[status]

    # 2. Income capacity score (scaled up to 35 points)
    income_points = min((income / 100_000) * 35, 35)

    # 3. Credit score rating (scaled against 850, up to 35 points)
```

```

credit_points = (credit_score / 850) * 35

# 4. Debt-to-Loan Affordability (15 points if loan amount is <= 2x annual income)
affordability_points = 15 if loan_amount <= income * 2 else 0

total_score = income_points + credit_points + employment_score + affordability_points

return {
    "decision": "Approved" if total_score >= APPROVAL_THRESHOLD else "Rejected",
    "score": round(total_score, 2),
    "factors": {
        "income_score": round(income_points, 2),
        "credit_score": round(credit_points, 2),
        "employment_status": employment_score,
        "loan_affordability": affordability_points,
    },
    "sensitive_attributes_used": False
}


def main():
    print("--- FAIR LOAN EVALUATION SYSTEM ---")
    income = float(input("Annual income: "))
    age = int(input("Age: "))
    employment_status = input("Employment status (full-time/part-time/self-employed/unemployed): ")
    credit_score = float(input("Credit score (300-850): "))
    loan_amount = float(input("Loan amount: "))

    result = loan_decision_fair(income, age, employment_status, credit_score, loan_amount)
    print(f"\nFinal Decision: {result['decision']} (Composite Score: {result['score']}/100)")
    print(f"Risk Factors Evaluated: {result['factors']}")

if __name__ == "__main__":
    main()

```

Execution Output (Fair Revised):



```

C:\Users\drona\Downloads\AI Code\lab5.py
• (project1) PS C:\Users\drona\Downloads\AI Code> & "c:\Users\drona\Downloads\AI Code\project1\Scripts\python
exe" "c:/Users/drona/Downloads/AI Code/lab5.py"
FAIRER LOAN APPROVAL DEMONSTRATION
Annual income: 10000
Age: 18
Employment status (full-time/part-time/self-employed/unemployed): unemployed
Credit score (0-850): 750
Loan amount: 9999
{'decision': 'Rejected', 'score': 55.88, 'factors': {'income': 10.0, 'credit_score': 30.88, 'employment_status': 0, 'loan_affordability': 15}, 'sensitive_attributes_used': False}
• (project1) PS C:\Users\drona\Downloads\AI Code>

```

Figure 6.2: Terminal output of objective loan evaluation based strictly on financial parameters.

6.3 Technical Comparison Table

Detailed comparison of parameters and legal compliance across both loan systems:

Assessment Factor	Biased Lending Engine	Fair Financial Risk Engine	Regulatory & Ethical Impact
Applicant Attributes Used	Gender, religion, race, age, income, credit	Income, credit score, employment stability, loan	Removes unconstitutional and discriminatory personal

		amount	criteria
Legal Compliance	Violates ECOA, Fair Housing Act & Art. 15	Fully compliant with international fair lending mandates	Eliminates substantial litigation and regulatory penalty risks
Parity for Equal Profiles	Identical financial applicants receive opposing decisions	Identical financial profiles receive identical evaluations	Guarantees procedural justice and merit-based credit allocation
Decision Transparency	Conceals arbitrary demographic deductions	Exposes exact financial factors and sub-scores	Allows applicants to review and repair their creditworthiness
Societal Impact	Exacerbates historical wealth disparities	Promotes financial inclusion and economic opportunity	Ensures equitable capital allocation across diverse communities

Student Reflection & Ethical Lending Takeaway (Task 6):

Designing responsible credit systems requires rigorous attribute auditing. By eliminating sensitive demographic inputs and basing decisions strictly on objective financial capacity and debt affordability, we create algorithms that are legally compliant, mathematically defensible, and socially fair.

Comprehensive Lab Synthesis & Responsible AI Framework

Through these six exercises, we have demonstrated that Responsible AI is not an abstract philosophical concept, but a tangible software engineering discipline. The table below summarizes the core pillars practiced throughout this lab:

Responsible AI Pillar	Initial Naive AI Risk	Human-Engineered Solution	Standard / Reference
1. Data Privacy & Security	Plaintext password storage and terminal echo	PBKDF2-HMAC-SHA256 with random salt + getpass masking	OWASP, NIST SP 800-63B
2. Algorithmic Fairness	Arbitrary location bonus creating proxy bias	Normalized merit/need scoring with audited demographic logs	Procedural Justice Frameworks
3. Model Transparency	Opaque black-box binary decisions with no explanation	Itemized factor weights, cutoff disclosure, and clear reason	EU AI Act 'Right to Explanation'
4. Ethical Metrics	Toxic presenteeism via 60% attendance weighting	Balanced deliverables (50%), teamwork (35%), leave protection	Objective Performance Rubrics
5. Inclusive Design	Exclusionary binary gender and invasive disability questions	Optional preferred identity, anonymity, barrier remediation	W3C WCAG 2.2 Principles
6. Anti-Discrimination	Protected attribute bias (gender, religion, race)	Pure financial risk modeling (income, credit score, DTI ratio)	Equal Credit Opportunity Act (ECOA)

Conclusion: When utilizing Generative AI assistants in software development, engineers must act as vigilant auditors rather than passive consumers. AI models naturally optimize for syntactically correct and concise code, but often overlook critical security safeguards, systemic biases, and edge-case inclusivity. Ensuring that algorithms remain secure, equitable, explainable, and accessible is the fundamental responsibility of the human engineer.